

Practical Machine Learning-Project

Santosh Pawar

November 27, 2017

Title: Quantifying exercise patterns and activities of individuals from accelerometer data

Executive summary:

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. They Participants were asked to perform barbell lifts correctly and incorrectly in 5 different ways. This project used data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants to predict the manner in which they did the exercise.

Data Preprocessing

```
#Read the Data
trainRaw <- read.csv("pml-training.csv")
testRaw <- read.csv("pml-testing.csv")

dim(trainRaw)
```

```
## [1] 19622 160
```

```
dim(testRaw)
```

```
## [1] 20 160
```

The training data set contains 19622 observations (rows) and 160 variables (columns), while the testing data set contains 20 observations and 160 variables. The “classe” variable in the training set is the outcome to predict.

Clean the data

```
#Data is cleaned to remove observations with missing values and some meaningless variables

sum(complete.cases(trainRaw))
```

```
## [1] 406
```

```
#Remove columns that contain NA missing values
```

```
trainRaw <- trainRaw[, colSums(is.na(trainRaw)) == 0]
testRaw <- testRaw[, colSums(is.na(testRaw)) == 0]
```

```
#Remove columns that do not contribute much to the accelerometer measurements
```

```
classe <- trainRaw$classe
trainRemove <- grepl("^X|timestamp|window", names(trainRaw))
trainRaw <- trainRaw[, !trainRemove]
trainCleaned <- trainRaw[, sapply(trainRaw, is.numeric)]
trainCleaned$classe <- classe

testRemove <- grepl("^X|timestamp|window", names(testRaw))
testRaw <- testRaw[, !testRemove]
testCleaned <- testRaw[, sapply(testRaw, is.numeric)]

dim(trainCleaned)
```

```
## [1] 19622    53
```

```
dim(testCleaned)
```

```
## [1] 20 53
```

The cleaned training data set now contains 19622 observations and 53 variables, while the testing data set contains 20 observations and 53 variables. The “classe” variable is still in the cleaned training set.

Slice the data

```
#Split the cleaned training set into a training data set (70%) and a validation data set (30%). The val
```

```
# Set seed for reproducibility of data
set.seed(12345)
```

```
#Partition data
```

```
inTrain <- createDataPartition(trainCleaned$classe, p=0.70, list=F)
trainData <- trainCleaned[inTrain, ]
testData <- trainCleaned[-inTrain, ]
```

Data Modeling

```
#Fit a predictive model for activity recognition using Random Forest algorithm. Five-fold cross validat
```

```
controlRf <- trainControl(method="cv", 5)
```

```
modelRf <- train(classe ~ ., data=trainData, method="rf", trControl=controlRf, ntree=250)
modelRf
```

```
## Random Forest
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10990, 10989, 10990, 10989, 10990
```

```
## Resampling results across tuning parameters:
##
##   mtry  Accuracy  Kappa
##    2    0.9900996 0.9874758
##   27    0.9898812 0.9871998
##   52    0.9840573 0.9798299
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

Validation

#Model performance is predicted on the validation data set.

```
predictRf <- predict(modelRf, testData)
confusionMatrix(testData$classe, predictRf)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1673    1    0    0    0
##           B   10 1125    4    0    0
##           C    0   14 1011    1    0
##           D    0    0   24  940    0
##           E    0    0    0    3 1079
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9903
##           95% CI   : (0.9875, 0.9927)
##           No Information Rate : 0.286
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##           Kappa : 0.9877
##           McNemar's Test P-Value : NA
##
```

```
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9941  0.9868  0.9731  0.9958  1.0000
## Specificity      0.9998  0.9970  0.9969  0.9951  0.9994
## Pos Pred Value   0.9994  0.9877  0.9854  0.9751  0.9972
## Neg Pred Value   0.9976  0.9968  0.9942  0.9992  1.0000
## Prevalence       0.2860  0.1937  0.1766  0.1604  0.1833
## Detection Rate   0.2843  0.1912  0.1718  0.1597  0.1833
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9969  0.9919  0.9850  0.9955  0.9997
```

Model accuracy

```
#Model accuracy is estimated
accuracy <- postResample(predictRf, testData$classe)
accuracy
```

```
## Accuracy      Kappa
## 0.9903144 0.9877458
```

Out of Sample Error estimation

```
#Estimate the out of sample error
oose <- 1 - as.numeric(confusionMatrix(testData$classe, predictRf)$overall[1])
oose
```

```
## [1] 0.009685641
```

The accuracy of the model is estimated to be 99.03% and the estimated out-of-sample error is 0.97%.

Predicting for Test Data Set


```
#Apply the training data model to the testing data set downloaded from the data source. We remove the p
result <- predict(modelRf, testCleaned[, -length(names(testCleaned))])
result
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Appendix

Figures

```
#Fig.1: Correlation Matrix Visualization
corrPlot <- cor(trainData[, -length(names(trainData))])
corrplot(corrPlot, method="color")
```



```
treeModel <- rpart(classe ~ ., data=trainData, method="class")
prp(treeModel)
```

